

DISTRIBUTED SYSTEMS ASSIGNMENT REPORT



Assignment ID: 4

Student Name: 徐春晖 XU Chunhui

Student ID: 12110304

Task 0: K8s Deployment & Service

Structure



```
1 |─codebase
2 | |─flask_app # app
3 | | | Makefile
4 | | | README.md
5 | | | Dockerfile # Dockerfile
6 | | | requirements.txt # pip requirement
7 | | |
8 | | |─v1.0.0 # old api version (without /chat)
9 | | |   app.py
10 | | |
11 | | |─v1.0.1 # new api version (have /chat)
12 | | |   app.py
13 | | |
14 | |─t0 # task 0
15 | |   kind-config.yaml
16 | |   Makefile # commands
17 | |   t0.yaml
18 | |
19 | |─t1 # task 1
20 | |   kind-config.yaml
21 | |   Makefile # commands
22 | |   t1-new.yaml # without class=vip tolerantion
23 | |   t1-old.yaml # with class=vip tolerantion
```

Build Image

Instructions

```
1 # codebase/flask_app/Makefile
2
3 .PHONY: build-old build-new build-all
4
5 OLD_VERSION:=1.0.0
6 NEW_VERSION:=1.0.1
7
8 build-old:
9     docker build --build-arg VERSION=$(OLD_VERSION) -t a4-flask:$(OLD_VERSION) .
10
11 build-new:
12     docker build --build-arg VERSION=$(NEW_VERSION) -t a4-flask:$(NEW_VERSION) .
13
14 build-all: build-old build-new
```

Screenshot



```
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\flask_app > main
> make build-all
docker build --build-arg VERSION=1.0.0 -t a4-flask:1.0.0 .
[+] Building 0.2s (9/9) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 291B
=> [internal] load metadata for docker.io/library/python:3.13.1-slim
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [1/4] FROM docker.io/library/python:3.13.1-slim
=> [internal] load build context
=> => transferring context: 711B
=> CACHED [2/4] WORKDIR /app
=> CACHED [3/4] COPY requirements.txt ./1.0.0/app.py ./
=> CACHED [4/4] RUN pip3 install --no-cache-dir -r requirements.txt -i https://pypi.tuna.tsinghua.edu.cn/simple
=> exporting to image
=> => exporting layers
=> => writing image sha256:0ec91cd641d729d347a7a03ce0de7e0d450749a72aa607fcf21ce998cbaf039f
=> => naming to docker.io/library/a4-flask:1.0.0

View build details: docker-desktop://dashboard/build/desktop-linux/desktop-linux/71wj1glvfgrizerqk9hx1ra4p
```

Task 0: K8s Deployment & Service

Instructions

I use these necessary command in task 0.

Actually, I'm using the original commands in the screenshot so that it's easier to tell.


```
1 # codebase/t0/Makefile
2
```

```

3  .PHONY: create-cluster delete-cluster apply-config update-image show-pods
4
5  CLUSTER_NAME := a4t0
6  DEPLOYMENT_NAME := a4t0
7  DEPLOYMENT_SERVICE := a4t0-service
8  OLD_IMAGE := a4-flask:1.0.0
9  NEW_IMAGE := a4-flask:1.0.1
10 K8S_CONFIG_FILE := t0.yaml
11
12 # Create a new cluster and load the images
13 create-cluster:
14     kind create cluster --name $(CLUSTER_NAME) --config kind-config.yaml
15     kind load docker-image $(OLD_IMAGE) --name $(CLUSTER_NAME)
16     kind load docker-image $(NEW_IMAGE) --name $(CLUSTER_NAME)
17
18 # Delete the cluster
19 delete-cluster:
20     kind delete cluster --name $(CLUSTER_NAME)
21
22 apply-config:
23     kubectl apply -f $(K8S_CONFIG_FILE)
24     kubectl set image deployment/$(DEPLOYMENT_NAME) flask-containers=$(OLD_IMAGE)
25
26 # Update the deployment with the new image
27 update-image:
28     kubectl set image deployment/$(DEPLOYMENT_NAME) flask-containers=$(NEW_IMAGE)
29     kubectl describe deployment $(DEPLOYMENT_NAME)
30
31 # Show the pods status
32 show-pods:
33     kubectl get pods -o wide

```

Cluster Configuration



```

1  # codebase/t0/kind-config.yaml
2
3  kind: cluster
4  apiVersion: kind.x-k8s.io/v1alpha4
5  nodes:
6  - role: control-plane
7  - role: worker
8  - role: worker
9  - role: worker

```

Result and Screenshot

Build cluster

```
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_2024f_Works\Assignment4\codebase\t0 > |main
> kind create cluster --name a4t0 --config kind-config.yaml
Creating cluster "a4t0" ...
  • Ensuring node image (kindest/node:v1.31.2)  0/1
  ✓ Ensuring node image (kindest/node:v1.31.2)  1/1
  • Preparing nodes 0/3
  ✓ Preparing nodes 3/3
  • Writing configuration 0/1
  ✓ Writing configuration 1/1
  • Starting control-plane 0/1
  ✓ Starting control-plane 1/1
  • Installing CNI 0/1
  ✓ Installing CNI 1/1
  • Installing StorageClass 0/1
  ✓ Installing StorageClass 1/1
  • Joining worker nodes 0/3
  ✓ Joining worker nodes 3/3
Set kubectl context to "kind-a4t0"
You can now use your cluster with:

kubectl cluster-info --context kind-a4t0

Have a nice day! 🌞
```

v1.0.0 Test and Result

```
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> kubectl get svc
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE
a4t0-service  ClusterIP   10.96.4.197    <none>         80/TCP     3m50s
kubernetes    ClusterIP   10.96.0.1      <none>         443/TCP    5m33s

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> docker exec -it a4t0-control-plane curl 10.96.4.197:80/
Hello! This is server in pod "<a4t0-559cbb74c8-p8d9x>" (IP=<10.244.1.3>) from node "<a4t0-worker3>"!
What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug a4t0-control-plane
  Learn more at https://docs.docker.com/go/debug-cli/

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> docker exec -it a4t0-control-plane curl 10.96.4.197:80/
Hello! This is server in pod "<a4t0-559cbb74c8-dfk96>" (IP=<10.244.3.2>) from node "<a4t0-worker>"!
What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug a4t0-control-plane
  Learn more at https://docs.docker.com/go/debug-cli/

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> docker exec -it a4t0-control-plane curl 10.96.4.197:80/chat/x
<!doctype html>
<html lang=en>
<title>404 Not Found</title>
<h1>Not Found</h1>
<p>The requested URL was not found on the server. If you entered the URL manually please check your spelling and try again.</p>
```

From first two commands, we can see the service in different node is running correctly. And we can see pod IP and node name from the `/` API response.

We can see that the access of `/chat` API will cause 404 error. The Flask controller will return a piece of 404 HTML code, indicated that this path is unavailable.

Load Image v1.0.1

```
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> kind load docker-image a4-flask:1.0.1 --name a4t0
Image: "a4-flask:1.0.1" with ID "sha256:9e7feb781edb16adc4e4f565c88993aacee59976e411f52cccfdc14b918a7ac5" found to be already present on all nodes

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> kubectl set image deployment/a4t0 flask-containers=a4-flask:1.0.1
deployment.apps/a4t0 image updated

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> kubectl rollout status deployment/a4t0
deployment "a4t0" successfully rolled out

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t0 > |main
> docker exec -it a4t0-control-plane curl 10.96.215.76:80/chat/octcarp?institution=sustech
{"message":"Hello octcarp from sustech!"}
```

We use new image, roll out and perform `/chat` API, its function is normal.

Rollout Events

```
OldReplicaSets: a4t0-559cbb74c8 (0/0 replicas created)
NewReplicaSet: a4t0-58d6965f55 (4/4 replicas created)
Events:
  Type    Reason             Age   From                  Message
  ----    -
  Normal  ScalingReplicaSet  37s   deployment-controller Scaled up replica set a4t0-559cbb74c8 to 4
  Normal  ScalingReplicaSet  17s   deployment-controller Scaled up replica set a4t0-58d6965f55 to 1
  Normal  ScalingReplicaSet  17s   deployment-controller Scaled down replica set a4t0-559cbb74c8 to 3 from 4
  Normal  ScalingReplicaSet  17s   deployment-controller Scaled up replica set a4t0-58d6965f55 to 2 from 1
  Normal  ScalingReplicaSet  16s   deployment-controller Scaled down replica set a4t0-559cbb74c8 to 2 from 3
  Normal  ScalingReplicaSet  16s   deployment-controller Scaled up replica set a4t0-58d6965f55 to 3 from 2
  Normal  ScalingReplicaSet  16s   deployment-controller Scaled down replica set a4t0-559cbb74c8 to 1 from 2
  Normal  ScalingReplicaSet  16s   deployment-controller Scaled up replica set a4t0-58d6965f55 to 4 from 3
  Normal  ScalingReplicaSet  15s   deployment-controller Scaled down replica set a4t0-559cbb74c8 to 0 from 1
```



```
1 # codebase/t0/t0.yaml
2
3 strategy:
4   rollingUpdate:
5     maxSurge: 1
6     maxUnavailable: 1
7   type: RollingUpdate
```

- `maxUnavailable: 1` : At most one pod can be unavailable during the update
- `maxSurge: 1` : At most one extra pod can be created during the update

And we can see that the new replica set plus 1 with the old replica set minus 1. Finally, all the 4 replicas becomes new.

This is my rollout strategy result.

Delete Pod

I ran the following command:



```
1 kubectl delete pod a4t0-559cbb74c8-2zt2x
```

Then I got:

```
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works > |main Carmen@Xpro 03:00:22
kubectll get pods -w
NAME                                READY   STATUS    RESTARTS   AGE
a4t0-559cbb74c8-2zt2x              1/1     Running   0          14m
a4t0-559cbb74c8-7cqdt             1/1     Running   0          14m
a4t0-559cbb74c8-kqdt              1/1     Running   0          14m
a4t0-559cbb74c8-tw7gs             1/1     Running   0          14m
a4t0-559cbb74c8-2zt2x             1/1     Terminating   0          15m
a4t0-559cbb74c8-jwcl7             0/1     Pending     0          0s
a4t0-559cbb74c8-jwcl7             0/1     Pending     0          0s
a4t0-559cbb74c8-jwcl7             0/1     ContainerCreating   0          0s
a4t0-559cbb74c8-2zt2x             0/1     Completed    0          15m
a4t0-559cbb74c8-2zt2x             0/1     Completed    0          15m
a4t0-559cbb74c8-2zt2x             0/1     Completed    0          15m
a4t0-559cbb74c8-jwcl7            1/1     Running      0          1s
```

From the screenshots, we can see that: the old pod is completed and a new one becomes running.

If a pod dies, K8s will automatically create a new pod.

It will always keep the number of running pods equal to the value specified by replicas

Task 1: K8s Pod Scheduling

Instructions

I use these necessary command in task 1.

Actually, I'm using the original commands in the screenshot so that it's easier to tell.

```
1  # codebase/t1/Makefile
2
3  .PHONY: create-cluster delete-cluster list-node-labels list-node-taints apply-old apply-
new
4  .PHONY: show-pods show-describe scale scale-1 scale-2 scale-3 scale-4 scale-5 show-pods
5
6  CLUSTER_NAME := a4t1
7  DEPLOYMENT_NAME := a4t1-deployment
8  K8S_CONFIG_FILE_OLD := t1-old.yaml
9  K8S_CONFIG_FILE_NEW := t1-new.yaml
10
11 create-cluster:
12     kind create cluster --name $(CLUSTER_NAME) --config kind-config.yaml
13
14 delete-cluster:
15     kind delete cluster --name $(CLUSTER_NAME)
16
17 list-node-labels:
18     # kubectl get nodes --show-labels
19     kubectl get nodes -o jsonpath='{range .items[*]}{.metadata.name}{"\t"}
{.metadata.labels}{"\n"}{end}'
20
21 list-node-taints:
22     kubectl get nodes -o jsonpath='{range .items[*]}{.metadata.name}{"\t"}{.spec.taints}
{"\n"}{end}'
23
24 apply-old:
25     kubectl apply -f $(K8S_CONFIG_FILE_OLD)
26
27 apply-new:
28     kubectl apply -f $(K8S_CONFIG_FILE_NEW)
29
30 show-describe:
31     kubectl describe deployment $(DEPLOYMENT_NAME)
32
33 show-pods:
34     kubectl get pods -o wide
35
36 scale:
37     kubectl scale deployment $(DEPLOYMENT_NAME) --replicas=$(REPLICAS)
38
39 scale-1: REPLICAS=1
40 scale-1: scale
```

```
41
42 scale-2: REPLICAS=2
43 scale-2: scale
44
45 scale-3: REPLICAS=3
46 scale-3: scale
47
48 scale-4: REPLICAS=4
49 scale-4: scale
50
51 scale-5: REPLICAS=5
52 scale-5: scale
```

Cluster Configuration



```
1 # codebase/t1/kind-config.yaml
2
3 kind: Cluster
4 apiVersion: kind.x-k8s.io/v1alpha4
5 nodes:
6 - role: control-plane
7 - role: worker
8   labels:
9     usage: normal
10 - role: worker
11   labels:
12     usage: normal
13     capability: powerful
14 - role: worker
15   kubeadmConfigPatches:
16   - |
17     kind: JoinConfiguration
18     nodeRegistration:
19       kubeletExtraArgs:
20         # no tier label
21         node-labels: "usage=normal,capability=powerful"
22       taints:
23         - key: class
24           value: vip
25           effect: NoSchedule
26 - role: worker
27   labels:
28     usage: backup
29 - role: worker
30   labels:
31     usage: backup
32
```

Result and Screenshot

Build cluster

```
E:\SUSTech\UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > main
> kind create cluster --name a4t1 --config kind-config.yaml
Creating cluster "a4t1" ...
  • Ensuring node image (kindest/node:v1.31.2)  ...
  ✓ Ensuring node image (kindest/node:v1.31.2)  ...
  • Preparing nodes  ...
  ✓ Preparing nodes  ...
  • Writing configuration  ...
  ✓ Writing configuration  ...
  • Starting control-plane  ...
  ✓ Starting control-plane  ...
  • Installing CNI  ...
  ✓ Installing CNI  ...
  • Installing StorageClass  ...
  ✓ Installing StorageClass  ...
  • Joining worker nodes  ...
  ✓ Joining worker nodes  ...
Set kubectl context to "kind-a4t1"
You can now use your cluster with:

kubectl cluster-info --context kind-a4t1
```

Scale without Toleration

```
1 # codebase/t1/t1_old.yaml
2
3 spec:
4   affinity:
5     podAntiAffinity:
6       requiredDuringSchedulingIgnoredDuringExecution:
7       - labelSelector:
8         matchExpressions:
9         - key: app
10           operator: In
11           values:
12           - a4t1
13         topologyKey: "kubernetes.io/hostname"
14     nodeAffinity:
15       preferredDuringSchedulingIgnoredDuringExecution:
16       - weight: 100
17         preference:
18         matchExpressions:
19         - key: capability
20           operator: In
21           values:
22           - powerful
23       - weight: 50
24         preference:
25         matchExpressions:
26         - key: usage
27           operator: NotIn
28           values:
29           - backup
```

In this file, we specify the scheduling rules

- Use `podAntiAffinity` to distribute multiple pods of the same service to different nodes to avoid single node failure.

- Use `nodeAffinity` to set the weight, let nodes with higher weights be scheduled with higher priority
 - If a node has label `capability: powerful` it will get 100 weight.
 - If a node doesn't have label `usage: backup` it will get 50 weight.
 - Overall, Worker 2, 3 have 150 weight, Worker has 50, Worker 4, 5 has 0.

```

> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl apply -f t1-old.yaml
deployment.apps/a4t1-deployment configured
service/a4t1-service unchanged
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl scale deployment a4t1-deployment --replicas=1
deployment.apps/a4t1-deployment scaled
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl scale deployment a4t1-deployment --replicas=2
deployment.apps/a4t1-deployment scaled
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl scale deployment a4t1-deployment --replicas=3
deployment.apps/a4t1-deployment scaled
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl scale deployment a4t1-deployment --replicas=4
deployment.apps/a4t1-deployment scaled
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl scale deployment a4t1-deployment --replicas=5
deployment.apps/a4t1-deployment scaled
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GAT
a4t1-deployment-7c7b777d49-8dtnf	1/1	Running	0	17s	10.244.5.7	a4t1-worker5	<none>		<none>	
a4t1-deployment-7c7b777d49-cqbdc	0/1	Pending	0	7s	<none>	<none>	<none>		<none>	
a4t1-deployment-7c7b777d49-ghcdp	1/1	Running	0	26s	10.244.1.6	a4t1-worker4	<none>		<none>	
a4t1-deployment-7c7b777d49-k557h	1/1	Running	0	47s	10.244.4.7	a4t1-worker2	<none>		<none>	
a4t1-deployment-7c7b777d49-wb54r	1/1	Running	0	34s	10.244.2.8	a4t1-worker	<none>		<none>	

We can see that, during the scale from 0 - 5, each Pod are in different node (due to Pod Anti-Affinity).

And from the `AGE`, we can see the order in which Pods are created in a Node is:

- Worker 2 -> 1 -> 4 -> 5 -> 3(unavailable).

This is consistent with our preset functional weights. The taint label of Worker 3 works normally.

```

> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl describe deployment a4t1-deployment
Name: a4t1-deployment
Namespace: default
CreationTimestamp: Wed, 18 Dec 2024 00:27:48 +0800
Labels: <none>
Annotations: deployment.kubernetes.io/revision: 4
Selector: app=a4t1
Replicas: 5 desired | 5 updated | 5 total | 4 available | 1 unavailable

```

Due to taint label, replica of Node Worker 3 is unavailable.

Scale with Toleration

Modify K8s configuration in `t1_new.yaml`, tolerant taints `class=vip`

```

1 # codebase/t1/t1_new.yaml
2
3     tolerations:
4     - key: "class"
5       operator: "Equal"
6       value: "vip"
7       effect: "NoSchedule"

```

```

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl apply -f t1-new.yaml
deployment.apps/a4t1-deployment configured
service/a4t1-service unchanged
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl scale deployment a4t1-deployment --replicas=1
deployment.apps/a4t1-deployment scaled
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl scale deployment a4t1-deployment --replicas=2
deployment.apps/a4t1-deployment scaled
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl scale deployment a4t1-deployment --replicas=3
deployment.apps/a4t1-deployment scaled
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl scale deployment a4t1-deployment --replicas=4
deployment.apps/a4t1-deployment scaled
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl scale deployment a4t1-deployment --replicas=5
deployment.apps/a4t1-deployment scaled
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GA
a4t1-deployment-745775d989-nv8gz	1/1	Running	0	27s	10.244.2.10	a4t1-worker	<none>		<none>	
a4t1-deployment-745775d989-q7c9b	1/1	Running	0	16s	10.244.1.8	a4t1-worker4	<none>		<none>	
a4t1-deployment-745775d989-qlrr2	1/1	Running	0	6s	10.244.5.9	a4t1-worker5	<none>		<none>	
a4t1-deployment-745775d989-r6qmf	1/1	Running	0	58s	10.244.4.9	a4t1-worker2	<none>		<none>	
a4t1-deployment-745775d989-vm924	1/1	Running	0	35s	10.244.3.3	a4t1-worker3	<none>		<none>	

When we tolerant taint with `class=vip`, we can use Node Worker 3 normally. And Its capability is `powerful`, so the scheduling order is:

- Worker 2 -> 3 -> 1 -> 4 -> 5.

```

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl describe deployment a4t1-deployment
Name: a4t1-deployment
Namespace: default
CreationTimestamp: Wed, 18 Dec 2024 00:27:48 +0800
Labels: <none>
Annotations: deployment.kubernetes.io/revision: 5
Selector: app=a4t1
Replicas: 5 desired | 5 updated | 5 total | 5 available | 0 unavailable

```

All the 5 replicas is available.

API Test

```

E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > kubectl get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
a4t1-service ClusterIP 10.96.103.130 <none> 80/TCP 34m
kubernetes ClusterIP 10.96.0.1 <none> 443/TCP 35m
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > docker exec -it a4t1-control-plane curl 10.96.103.130:80/
Hello! This is server in pod "<a4t1-deployment-745775d989-nv8gz>" (IP=<10.244.2.10>) from node "<a4t1-worker>"!
What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image -> docker debug a4t1-control-plane
Learn more at https://docs.docker.com/go/debug-cli/
E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > docker exec -it a4t1-control-plane curl 10.96.103.130:80/chat/octcarp?institution=sustech
{"message":"Hello octcarp from sustech!"}

```

The function is normal.

Problems

Image pull problem

In task 1, I specify the container's image instead of manually load it into pods.

When I specify the image as local `a4-flask:1.0.1` in `t1.yaml`, an `ErrImagePull` error will occur. And the pod will not working correctly.

```
> E:\SUSTech_UGCS\SUSTech_CS328-Distributed_F24_Works\Assignment4\codebase\t1 > |main
> kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE           NOMINATED NODE   READINE
SS GATES
a4t1-deployment-7fc4c59d7d-6lq2l    0/1     ErrImagePull  0          14s   10.244.4.11   a4t1-worker2   <none>           <none>
```

The reason is that Docker will first look for the corresponding image from DockerHub, and if it is not found, it will report an error.

After searching, I also found that creating a Docker local registry can also solve this problem. However, for convenience, I still re-tagged the local image, uploaded it to my personal DockerHub domain, and specified the image as a cloud image to solve it.

```
1 # codebase/t1/t1_new.yaml
2
3 # old version
4 containers:
5   - name: flask-container
6     image: a4-flask:1.0.1
7
8 # new version
9 containers:
10  - name: flask-container
11    image: octcarp/sustech-cs328:a4-flask-new
12    imagePullPolicy: IfNotPresent
```

Bonus (*): Advice on Future Cloud Computing Lab

In Assignment 1

Overall, the experience was great. The implementation of MPI parallel computing in C made me feel the charm of distributed computing.

Advice

It is possible to explicitly require to implement several different forms of MPI (point-to-point communication or broadcast), and on this basis compare the operating efficiency of different numbers of processes, and make horizontal and vertical comparisons.

In addition to matrix multiplication, parallel computing scenarios can be expanded, such as parallel sorting algorithms.

In Assignment 2

The part I enjoy

- Implements cross-language microservice modules, especially Go-based ones. This made me realize the unity and efficiency of the gRPC protocol more deeply.
- Also, it was really fun to make several microservice modules and make them work together in a docker compose network.

The part I struggle with

- Design the service logic of RESTful API. Although the service logic is not particularly complex, but design it is quite tiring. This part is covered in many other Web application design courses. I think we could focus on the deployment of distributed microservices. Maybe TA could provide some API design demo first?
- The unknown port occupation of the Kafka service may be a problem (but this may be my own problem).

Advice

In general, it allows us to implement more complex distributed/microservice architectures, instead of spending too much effort on API design.

The demo architecture of the assignment can be more detailed. I spent a lot of time thinking about how to organize the file structure.

In Assignment 3

Advice

I was a little constrained by the fact that I could only use `Pyspark`, and I wanted to try using `Scala` to complete this assignment. Perhaps the language limitation could be relaxed in later semesters.

Maybe some `MapReduce` programming could be involved in the assignment.

In Assignment 4

It was a good experience. I became more familiar with the operating principles of `K8s` during the experiment. Understand the affinity / anti-affinity of Pod / Node, and the use of taints and toleration.

Advice

Could introduce the use of other tools such as `k3s`.

More Topics

- CUDA programming? I'm not sure if this is easy to implement (because NVIDIA GPU are required), but there will be many scenarios involving CUDA programming in the future, so it feels good to learn about it.
- A more detailed load balancing design experiment: Use different load balancing algorithms to compare their different focuses and overall effects.